

Survey Analysis

Zachary Yeap

```
library(tidyverse)
library(janitor)
library(knitr)

# Update the filename if needed
survey <- read_csv("Survey (Responses) - Form responses 1.csv")

# Total responses + column count
dim(survey)

# Column names (so we can confirm structure matches last round)
colnames(survey)

# Peek at how Likert values are stored in the new data
# (Same as before we'd expect: "Strongly Disagree" ... "Strongly Agree" as text)
unique(survey[[7]])    # first Likert column
unique(survey[[27]])   # satisfaction column

# How many responses look usable at first glance?
glimpse(survey)

survey <- survey %>%
  rename(
    timestamp = 1,
    consent   = 2,
    age       = 3,
    gender     = 4,
    app_usage = 5,
    role      = 6,

    # Section B - Clarity
    clear_llm      = 7, clear_rule      = 8,
    unambiguous_llm = 9, unambiguous_rule = 10,
    easy_to_answer_llm = 11, easy_to_answer_rule = 12,
    oe_clarity     = 13,

    # Section C - Relevance
    relevant_persona_llm = 14, relevant_persona_rule = 15,
    addressed_priority_llm = 16, addressed_priority_rule = 17,
    reflected_change_llm  = 18, reflected_change_rule  = 19,

    # Section D - Quality
    captured_detail_llm = 20, captured_detail_rule = 21,
    well_designed_llm   = 22, well_designed_rule   = 23,
```

```

natural_flow_llm      = 24, natural_flow_rule      = 25,
oe_compare            = 26,

# Section E
satisfaction = 27,
oe_overall   = 28
)

```

```

recode_likert <- function(x) {
  case_when(
    x == "Strongly Disagree" ~ 1,
    x == "Disagree"          ~ 2,
    x == "Neutral"           ~ 3,
    x == "Agree"             ~ 4,
    x == "Strongly Agree"    ~ 5,
    TRUE                     ~ NA_real_
  )
}

likert_cols <- c(
  "clear_llm", "clear_rule",
  "unambiguous_llm", "unambiguous_rule",
  "easy_to_answer_llm", "easy_to_answer_rule",
  "relevant_persona_llm", "relevant_persona_rule",
  "addressed_priority_llm", "addressed_priority_rule",
  "reflected_change_llm", "reflected_change_rule",
  "captured_detail_llm", "captured_detail_rule",
  "well_designed_llm", "well_designed_rule",
  "natural_flow_llm", "natural_flow_rule"
)

survey <- survey %>% mutate(across(all_of(likert_cols), recode_likert))

```

```

n_raw <- nrow(survey)

# Filter 1: Drop non-consenting responses
survey_clean <- survey %>% filter(consent == "Yes")
n_after_consent <- nrow(survey_clean)

# Filter 2: Drop the troll response
# Identify by the specific invalid pattern: all 18 Likert items at 1 (Strongly Disagree)
# AND non-substantive open-ended responses ("Me", "You", "Mama")
survey_clean <- survey_clean %>%
  rowwise() %>%
  mutate(all_strongly_disagree = all(c_across(all_of(likert_cols)) == 1, na.rm = TRUE)) %>%
  ungroup() %>%
  filter(!all_strongly_disagree) %>%
  select(-all_strongly_disagree)

n_final <- nrow(survey_clean)

cat("--- Sample size chain ---\n")

```

```
## --- Sample size chain ---
```

```
cat("Raw responses:           ", n_raw, "\n")
```

```
## Raw responses:           42
```

```
cat("After consent filter:    ", n_after_consent, "\n")
```

```
## After consent filter:    41
```

```
cat("After troll removal:     ", n_final, "\n")
```

```
## After troll removal:     40
```

```
cat("Total excluded:          ", n_raw - n_final, "\n")
```

```
## Total excluded:          2
```

```
survey <- survey_clean
```

```
age_tab    <- survey %>% tabyl(age)           %>% adorn_pct_formatting(digits = 1)
gender_tab  <- survey %>% tabyl(gender)        %>% adorn_pct_formatting(digits = 1)
app_tab     <- survey %>% tabyl(app_usage)     %>% adorn_pct_formatting(digits = 1)
role_tab    <- survey %>% tabyl(role)          %>% adorn_pct_formatting(digits = 1)
```

```
kable(age_tab,      caption = "Age range")
```

Table 1: Age range

age	n	percent
18-20	8	20.0%
21-25	26	65.0%
26-34	6	15.0%

```
kable(gender_tab,  caption = "Gender")
```

Table 2: Gender

gender	n	percent
Female	21	52.5%
Male	19	47.5%

```
kable(app_tab,      caption = "Health app usage")
```

Table 3: Health app usage

app_usage	n	percent
Every day	3	7.5%
Never	14	35.0%
Never, Prefer not to say	1	2.5%
Once a week	5	12.5%
Rarely	15	37.5%
Twice a week	2	5.0%

```
kable(role_tab, caption = "Student / working")
```

Table 4: Student / working

role	n	percent
Student	28	70.0%
Working	12	30.0%

```
likert_summary_table <- function(data, item_stems) {
  data %>%
    select(starts_with(item_stems)) %>%
    pivot_longer(everything(),
                  names_to = c("item", "method"),
                  names_pattern = "(.+)_(llm|rule)",
                  values_to = "score") %>%
    group_by(item, method) %>%
    summarise(n      = sum(!is.na(score)),
              mean    = round(mean(score, na.rm = TRUE), 2),
              sd      = round(sd(score, na.rm = TRUE), 2),
              median   = median(score, na.rm = TRUE),
              .groups = "drop") %>%
    pivot_wider(names_from = method,
                values_from = c(n, mean, sd, median)) %>%
    mutate(diff_llm_minus_rule = round(mean_llm - mean_rule, 2)) %>%
    select(item,
           n_rule, mean_rule, sd_rule, median_rule,
           n_llm, mean_llm, sd_llm, median_llm,
           diff_llm_minus_rule)
}

likert_plot <- function(data, item_stems, plot_title) {
  data %>%
    select(starts_with(item_stems)) %>%
    pivot_longer(everything(),
                  names_to = c("item", "method"),
                  names_pattern = "(.+)_(llm|rule)",
                  values_to = "score") %>%
    group_by(item, method) %>%
    summarise(mean = mean(score, na.rm = TRUE), .groups = "drop") %>%
    ggplot(aes(x = item, y = mean, fill = method)) +
    geom_col(position = "dodge", width = 0.7) +

```

```

    geom_text(aes(label = round(mean, 2)),
              position = position_dodge(width = 0.7),
              vjust = -0.3, size = 3.5) +
    scale_y_continuous(limits = c(0, 5), breaks = 0:5) +
    labs(title = plot_title, x = NULL,
          y = "Mean Likert score (1-5)", fill = "Method") +
    theme_minimal()
}

likert_colnames <- c("Item",
                    "n (Rule)", "M (Rule)", "SD (Rule)", "Mdn (Rule)",
                    "n (LLM)", "M (LLM)", "SD (LLM)", "Mdn (LLM)",
                    "Diff (LLM-Rule)")

```

```

clarity_stems <- c("clear", "unambiguous", "easy_to_answer")
kable(likert_summary_table(survey, clarity_stems),
      caption = "Section B - Clarity: Rule-based vs LLM",
      col.names = likert_colnames)

```


Table 5: Section B — Clarity: Rule-based vs LLM

Item	n (Rule)	M (Rule)	SD (Rule)	Mdn (Rule)	n (LLM)	M (LLM)	SD (LLM)	Mdn (LLM)	Diff (LLM-Rule)
clear	40	4.03	0.95	4	40	4.22	0.73	4	0.19
easy_to_answer	40	4.03	1.05	4	40	4.10	0.96	4	0.07
unambiguous	40	3.70	1.07	4	40	3.83	0.96	4	0.13

```

likert_plot(survey, clarity_stems, "Clarity: Rule-based vs LLM")

```




```
relevance_stems <- c("relevant_persona", "addressed_priority", "reflected_change")
kable(likert_summary_table(survey, relevance_stems),
      caption = "Section C - Relevance: Rule-based vs LLM",
      col.names = likert_colnames)
```

Table 6: Section C — Relevance: Rule-based vs LLM

Item	n (Rule)	M (Rule)	SD (Rule)	Mdn (Rule)	n (LLM)	M (LLM)	SD (LLM)	Mdn (LLM)	Diff (LLM-Rule)
addressed_priority	40	3.62	0.87	4.0	40	4.00	0.75	4	0.38
reflected_change	40	3.60	0.93	3.5	40	4.10	0.78	4	0.50
relevant_persona	40	3.78	0.80	4.0	40	4.15	0.66	4	0.37

```
likert_plot(survey, relevance_stems, "Relevance: Rule-based vs LLM")
```




```
quality_stems <- c("captured_detail", "well_designed", "natural_flow")
kable(likert_summary_table(survey, quality_stems),
      caption = "Section D - Quality: Rule-based vs LLM",
      col.names = likert_colnames)
```

Table 7: Section D — Quality: Rule-based vs LLM

Item	n (Rule)	M (Rule)	SD (Rule)	Mdn (Rule)	n (LLM)	M (LLM)	SD (LLM)	Mdn (LLM)	Diff (LLM-Rule)
captured_detail	40	3.50	0.88	3.5	40	3.90	0.74	4	0.40
natural_flow	40	3.92	0.80	4.0	40	3.98	0.89	4	0.06
well_designed	40	3.58	0.93	4.0	40	4.00	0.78	4	0.42

```
likert_plot(survey, quality_stems, "Quality: Rule-based vs LLM")
```

Quality: Rule-based vs LLM



Per-participant composite scores: mean of 3 items per construct per method

```
survey <- survey %>%
  rowwise() %>%
  mutate(
    clarity_llm = mean(c(clear_llm, unambiguous_llm, easy_to_answer_llm), na.rm = TRUE),
    clarity_rule = mean(c(clear_rule, unambiguous_rule, easy_to_answer_rule), na.rm = TRUE),
    relevance_llm = mean(c(relevant_persona_llm, addressed_priority_llm, reflected_change_llm), na.rm = TRUE),
    relevance_rule = mean(c(relevant_persona_rule, addressed_priority_rule, reflected_change_rule), na.rm = TRUE),
    quality_llm = mean(c(captured_detail_llm, well_designed_llm, natural_flow_llm), na.rm = TRUE),
    quality_rule = mean(c(captured_detail_rule, well_designed_rule, natural_flow_rule), na.rm = TRUE)
  ) %>%
  ungroup()
```

```
paired_t_summary <- function(llm, rule, label) {
  tt <- t.test(llm, rule, paired = TRUE)
  diff <- llm - rule
  d <- mean(diff, na.rm = TRUE) / sd(diff, na.rm = TRUE)
  tibble(
    Construct = label,
    `M (LLM)` = round(mean(llm, na.rm = TRUE), 2),
    `M (Rule)` = round(mean(rule, na.rm = TRUE), 2),
    `Mean diff` = round(mean(diff, na.rm = TRUE), 2),
    t = round(tt$statistic, 2),
    df = tt$parameter,
    `p-value` = format.pval(tt$p.value, digits = 3, eps = .001),
    `Cohen's d` = round(d, 2)
  )
}
```



```

)
}

t_results <- bind_rows(
  paired_t_summary(survey$clarity_llm, survey$clarity_rule, "Clarity"),
  paired_t_summary(survey$relevance_llm, survey$relevance_rule, "Relevance"),
  paired_t_summary(survey$quality_llm, survey$quality_rule, "Quality")
)

kable(t_results, caption = "Paired t-tests: LLM vs Rule-based composite scores (n = 40)")

```

Table 8: Paired t-tests: LLM vs Rule-based composite scores (n = 40)

Construct	M (LLM)	M (Rule)	Mean diff	t	df	p-value	Cohen's d
Clarity	4.05	3.92	0.13	2.58	39	0.0139	0.41
Relevance	4.08	3.67	0.42	3.69	39	<0.001	0.58
Quality	3.96	3.67	0.29	3.76	39	<0.001	0.59

```

n_satisfaction <- sum(!is.na(survey$satisfaction))

satisfaction_endorsements <- survey %>%
  filter(!is.na(satisfaction)) %>%
  separate_longer_delim(satisfaction, delim = ", ") %>%
  count(satisfaction, name = "n") %>%
  mutate(percent = round(100 * n / n_satisfaction, 1)) %>%
  arrange(desc(n))

kable(satisfaction_endorsements,
  caption = paste0("Section E - Satisfaction endorsements (multi-select; % of n = ",
    n_satisfaction, " valid responses)"))

```

Table 9: Section E — Satisfaction endorsements (multi-select; % of n = 40 valid responses)


satisfaction	n	percent
Satisfied with both	22	55.0
Dynamic LLM-based method	17	42.5
Static Rule-based method	8	20.0
Neither	2	5.0

```

satisfaction_endorsements %>%
  mutate(satisfaction = fct_reorder(satisfaction, n)) %>%
  ggplot(aes(x = satisfaction, y = n, fill = satisfaction)) +
  geom_col(width = 0.6, show.legend = FALSE) +
  geom_text(aes(label = paste0(n, " (", percent, "%)")),
    hjust = -0.1, size = 4) +
  coord_flip() +
  scale_y_continuous(expand = expansion(mult = c(0, 0.2))) +

```

```
labs(title = "Which method(s) were participants satisfied with?",
     subtitle = "Multi-select - participants could endorse more than one option",
     x = NULL, y = "Count") +
theme_minimal()
```



```
oe_responses <- list(
  clarity = survey %>%
    filter(!is.na(oe_clarity), str_trim(oe_clarity) != "") %>%
    pull(oe_clarity),
  compare = survey %>%
    filter(!is.na(oe_compare), str_trim(oe_compare) != "") %>%
    pull(oe_compare),
  overall = survey %>%
    filter(!is.na(oe_overall), str_trim(oe_overall) != "") %>%
    pull(oe_overall)
)

# How many responses per question?
sapply(oe_responses, length)

cat("\n==== Q1: Confusing/unclear questions? =====\n\n")
walk(oe_responses$clarity, ~ cat("•", .x, "\n\n"))

cat("\n==== Q2: How did each version feel? Which preferred & why? =====\n\n")
walk(oe_responses$compare, ~ cat("•", .x, "\n\n"))
```

```
cat("\n==== Q3: Thoughts on the two methods - what worked / could improve? ==== \n\n")
walk(oe_responses$overall, ~ cat(".", .x, "\n\n"))
```

```
q1_themes <- tribble(
  ~Code, ~Count, ~Percent, ~`Example Response`,
  "No confusion / questions were clear", 34, "89.5%", "\"All questions were clear\"",
  "Timeframe ambiguity", 1, "2.6%", "\"confused how long I should check\"",
  "Need clearer instructions / onboarding", 1, "2.6%", "\"an initial guided tour for new users v",
  "Insufficient personalisation in both versions", 1, "2.6%", "\"the question are not personalise enough",
  "LLM questions felt too lengthy", 1, "2.6%", "\"LLM based questions were lengthier\"",
  "Positive: LLM examples were helpful", 1, "2.6%", "\"the llm helped by providing examples\"",
  "Non-substantive response", 1, "2.6%", "\".\""
)
kable(q1_themes, caption = "Q1 - Confusing or unclear questions? (n = 38)")
```

Table 10: Q1 — Confusing or unclear questions? (n = 38)

Code	Count	Percent	Example Response
No confusion / questions were clear	34	89.5%	“All questions were clear”
Timeframe ambiguity	1	2.6%	“confused how long I should check”
Need clearer instructions / onboarding	1	2.6%	“an initial guided tour for new users would be better”
Insufficient personalisation in both versions	1	2.6%	“the question are not personalise enough in both set”
LLM questions felt too lengthy	1	2.6%	“LLM based questions were lengthier”
Positive: LLM examples were helpful	1	2.6%	“the llm helped by providing examples”
Non-substantive response	1	2.6%	“.”

```
q2_themes <- tribble(
  ~Code, ~Count, ~Percent, ~`Example Response`,
  "Preferred LLM (overall)", 20, "51.3%", "\"prefer the llm\"",
  " - Reason: more personalised / tailored", 6, "15.4%", "\"felt more specific and better tai",
  " - Reason: adapts to history / changes", 4, "10.3%", "\"linked back to my month 1 question",
  " - Reason: easier / shorter / less effort", 3, "7.7%", "\"I like the LLM because it is shor",
  " - Reason: more natural / human-like", 2, "5.1%", "\"LLM feels more human-like\"",
  " - No specific reason stated", 5, "12.8%", "\"LLM based\"",
  "Preferred Rule-based", 4, "10.3%", "\"I prefer Rule-based because it is",
  "Felt both methods were similar", 1, "2.6%", "\"I felt they were quite similar\"",
  "Vague positive (no preference stated)", 8, "20.5%", "\"easy to answer and straightforward",
  "Non-substantive", 6, "15.4%", "\"N/A\" / \".\""
)
kable(q2_themes, caption = "Q2 - Preferred method and why? (n = 39)")
```

Table 11: Q2 — Preferred method and why? (n = 39)

Code	Count	Percent	Example Response
Preferred LLM (overall)	20	51.3%	“prefer the llm”
— Reason: more personalised / tailored	6	15.4%	“felt more specific and better tailored to my personal situation”

Code	Count	Percent	Example Response
— Reason: adapts to history / changes	4	10.3%	“linked back to my month 1 questions”
— Reason: easier / shorter / less effort	3	7.7%	“I like the LLM because it is shorter”
— Reason: more natural / human-like	2	5.1%	“LLM feels more human-like”
— No specific reason stated	5	12.8%	“LLM based”
Preferred Rule-based	4	10.3%	“I prefer Rule-based because it is straightforward”
Felt both methods were similar	1	2.6%	“I felt they were quite similar”
Vague positive (no preference stated)	8	20.5%	“easy to answer and straightforward”
Non-substantive	6	15.4%	“N/A” / “.”

```

q3_themes <- tribble(
  ~Code, ~Count, ~Percent, ~`Example Response`,
  "Need more personalisation / depth / coverage", 4, "30.8%", "\"more personalisation in the LLM-based method would be beneficial\"",
  "Scoring / display bugs reported", 2, "15.4%", "\"Scores seemed to be inverted\"",
  "Suggested new features (open text, method choice)", 2, "15.4%", "\"short answer to help input other/new updates\"",
  "UI / navigation issues", 1, "7.7%", "\"side by side page was a bit hard to navigate\"",
  "LLM adaptivity / context-awareness praised", 1, "7.7%", "\"adaptive nature of the LLM worked well\"",
  "Rule-based provided better context (alt view)", 1, "7.7%", "\"rule based gave more context\"",
  "Need clearer instructions / guidance", 1, "7.7%", "\"more instructions would make the process easier\"",
  "Both methods seemed fine", 1, "7.7%", "\"both version look fine for me\""
)
kable(q3_themes, caption = "Q3 - What worked / could improve? (n = 13)")

```

Table 12: Q3 — What worked / could improve? (n = 13)

Code	Count	Percent	Example Response
Need more personalisation / depth / coverage	4	30.8%	“more personalisation in the LLM-based method would be beneficial”
Scoring / display bugs reported	2	15.4%	“Scores seemed to be inverted”
Suggested new features (open text, method choice)	2	15.4%	“short answer to help input other/new updates”
UI / navigation issues	1	7.7%	“side by side page was a bit hard to navigate”
LLM adaptivity / context-awareness praised	1	7.7%	“adaptive nature of the LLM worked well”
Rule-based provided better context (alt view)	1	7.7%	“rule based gave more context”
Need clearer instructions / guidance	1	7.7%	“more instructions would make the process easier”
Both methods seemed fine	1	7.7%	“both version look fine for me”

Results

Sample

Forty-two participants completed the survey. Two responses were excluded from analysis: one participant who did not consent (n = 1), and one whose response showed invalid patterning — rating all 18 Likert items

as *Strongly disagree* and providing non-substantive open-ended answers (“Me”, “You”, “Mama”). Final analytic sample: $n = 40$.

The sample was predominantly young (65.0% aged 21–25, 20.0% aged 18–20, 15.0% aged 26–34) and student-dominated (70.0% students, 30.0% working professionals). Gender was roughly balanced (52.5% female, 47.5% male). Health-app usage was generally low, with 72.5% reporting they used such apps rarely or never. All Likert items used a 5-point scale (1 = *Strongly disagree*, 5 = *Strongly agree*).

Clarity

Both versions were rated positively, with all means above 3.7. The LLM-based version scored marginally higher on each item: “clear and easy to understand” ($M = 4.22$ vs 4.03), “easy to know how to answer” ($M = 4.10$ vs 4.03), and “specific and unambiguous wording” ($M = 3.83$ vs 3.70). Differences ranged from $+0.07$ to $+0.19$. A paired-samples t-test on composite Clarity scores confirmed that the LLM-based version was rated significantly higher than the Rule-based version, $t(39) = 2.58$, $p = .014$, $d = 0.41$ — a small-to-medium effect. Wording specificity was the weakest item for both methods, suggesting a shared design challenge.

The corresponding open-ended question on confusing or unclear items ($n = 38$) supported the quantitative picture (Table 10). 89.5% ($n = 34$) of respondents reported no confusion (e.g. “*All questions were clear*”). The few substantive concerns each appeared in only one response: ambiguity about the recall window (“*confused how long I should check*”), a need for clearer onboarding instructions, a sense that LLM questions were too lengthy, and insufficient personalisation in both versions. One participant offered a positive observation specifically about the LLM, noting it “*helped by providing examples*” for one section.

Relevance

Relevance produced the largest method differences. The LLM scored meaningfully higher on all three items: “relevant to Sophie’s health profile” ($M = 4.15$ vs 3.78 , $+0.37$), “addressed what seemed most important” ($M = 4.00$ vs 3.62 , $+0.38$), and “reflected changes in Sophie’s condition” ($M = 4.10$ vs 3.60 , $+0.50$). The Rule-based median on the change-tracking item was 3.5, while the LLM median reached 4. A paired t-test on composite Relevance scores showed a highly significant difference favouring the LLM-based version, $t(39) = 3.69$, $p < .001$, $d = 0.58$ — a medium effect identifying contextual adaptivity as the LLM’s strongest comparative advantage.

Quality

Quality scores favoured the LLM on two of three items: “captured important details” ($M = 3.90$ vs 3.50 , $+0.40$) and “well designed for meaningful assessment” ($M = 4.00$ vs 3.58 , $+0.42$). However, “felt natural and appropriate” was essentially tied ($M = 3.98$ vs 3.92 , $+0.06$). A paired t-test on composite Quality scores showed a highly significant LLM advantage, $t(39) = 3.76$, $p < .001$, $d = 0.59$ — a medium effect. Participants did not experience the LLM as more pleasant to complete, but they did experience it as better-aimed at the persona’s situation — an advantage of fit rather than feel.

Overall experience

The multi-select satisfaction question showed broad acceptance of both methods, with the LLM endorsed more than twice as often as the Rule-based version: 55.0% ($n = 22$) selected “satisfied with both,” 42.5% ($n = 17$) the LLM-based version, 20.0% ($n = 8$) the Rule-based version, and 5.0% ($n = 2$) neither.

The open-ended preference question ($n = 39$) was thematically analysed using inductive coding (Table 11). A majority of respondents (51.3%, $n = 20$) explicitly preferred the LLM-based version, compared to 10.3% ($n = 4$) who preferred the Rule-based version and 2.6% ($n = 1$) who felt the two were similar. The remaining responses were vague positives without a stated preference (20.5%, $n = 8$) or non-substantive (15.4%, $n =$

6). Among those who explained their LLM preference, the most cited reasons were: more personalised or tailored to the persona (15.4%, $n = 6$; e.g. *“felt more specific and better tailored to my personal situation”*); adapted to history and changes over time (10.3%, $n = 4$; e.g. *“linked back to my month 1 questions”*); easier, shorter, or required less effort (7.7%, $n = 3$); and more natural or human-like (5.1%, $n = 2$). Five LLM-preferring participants (12.8%) stated their preference without elaboration. Four participants offered a counter-perspective, preferring the Rule-based version for its straightforwardness (*“I prefer Rule-based because it is straightforward”*) or because it incorporated the previous month’s state more transparently.

Summary

Across nine evaluation items spanning Clarity, Relevance, and Quality, the LLM-based questionnaire was rated equal to or better than the Rule-based version in every case, with all three composite differences statistically significant. The largest effects appeared in Relevance ($d = 0.58$) and Quality ($d = 0.59$), with a smaller effect for Clarity ($d = 0.41$), indicating that both methods can produce understandable questions but only the LLM-based method adapts meaningfully to an evolving persona. The qualitative data converged on the same conclusion: a 5:1 explicit preference for the LLM, with personalisation and history awareness as the most cited advantages. The near-tie on “natural flow” and the persistent demand for deeper personalisation show that the LLM’s current advantage is one of contextual fit rather than user experience, with clear room to grow.